

## Build and run applications in a Dockerless Kubernetes world

Jorge Morales  
OpenShift Developer Advocate  
DevConf India 2018  
August 5th - 11:45 - Room 1

# Me (aka Jorge Morales)



- Spanish by nature and by language
- Work at Red Hat
- OpenShift Developer Advocate
- Mostly Java developer
- Obsessed with improving the developer experience



<http://jorgemoral.es>



[@jorgemoralespou](https://twitter.com/jorgemoralespou)



[github.com/jorgemoralespou](https://github.com/jorgemoralespou)

A photograph of a dirt road winding through a forest. The road is light-colored and leads towards a bright opening in the trees. The sides of the road are covered in tall, dry grass and some evergreen trees. The overall tone is warm and slightly hazy.


Each day is a little  
bit of history.

José Saramago




# ONCE

UPON A TIME



# Linux Containers



**Kernel namespaces:** sandboxing processes from one another

**Control Groups (cgroups):** control process resource allocations

**Security:** capabilities drop (seccomp), Mandatory access control (SELinux, Apparmor)

# Containers

OS virtualization

shared OS

Same OS as host

Start small, add later



# VM's

Hardware virtualization

Dedicated OS

Independent OS

Start full, strip out later

VMs? Containers?  
All I want to do is program!  
Jeesh.





Making things easy is hard.

Ted Nelson





*Because I'm  
happy!*





With scale came complexity

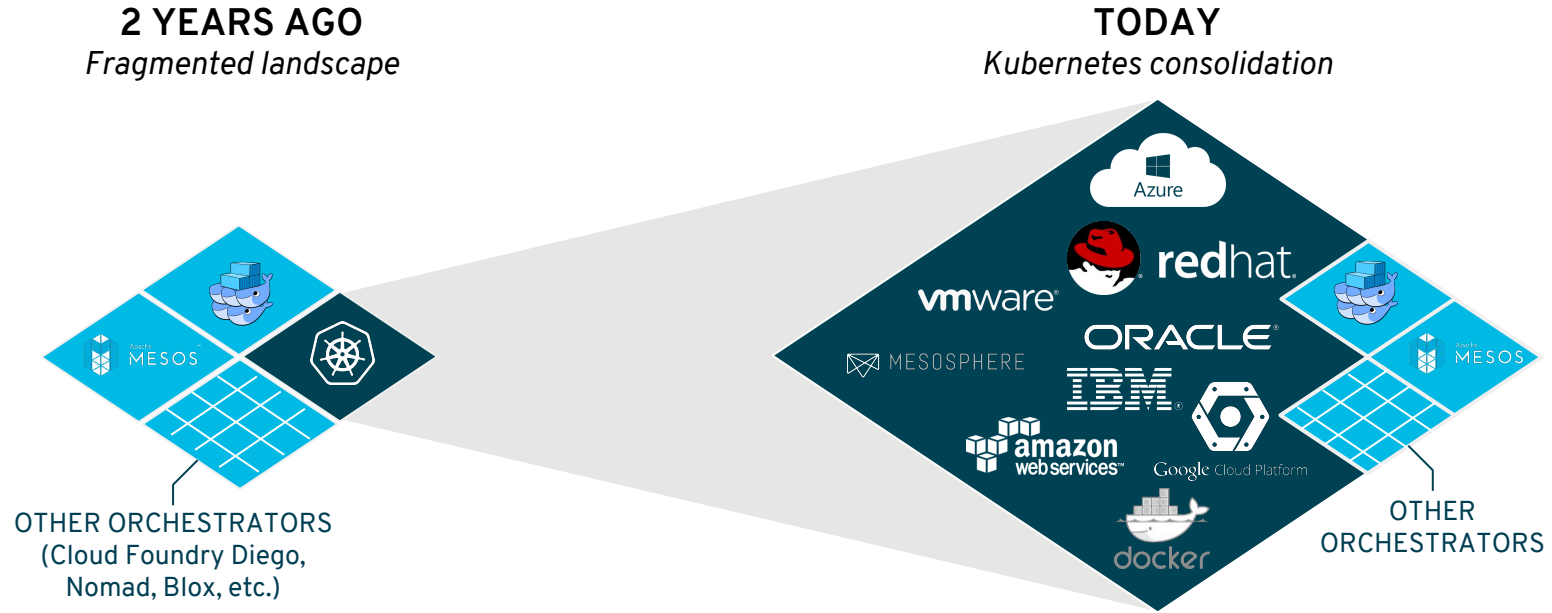


# Orchestration





# CONTAINER ORCHESTRATION LANDSCAPE



# Why kubernetes?



#1: Open source, backed by giants

#2: Vibrant and fast growing community

#3: Supported on all clouds

#4: Great partnerships

# Started slow

Docker 1.0

Kubernetes **1.0**: Supports  
Docker containers

June  
2014

Dec  
2014

July  
2015

Rkt 0.1.0

# then more runtimes showed up

Docker 1.0

Kubernetes **1.0**: Supports  
Docker containers

Kubernetes **1.3**: Supports  
Docker and Rkt containers

June  
2014

Dec  
2014

July  
2015

Feb  
2016


July  
2016

Rkt 0.1.0

Rkt 1.0



and code got messy



“Change is the essential  
process of all of  
existence.”

—SPOCK

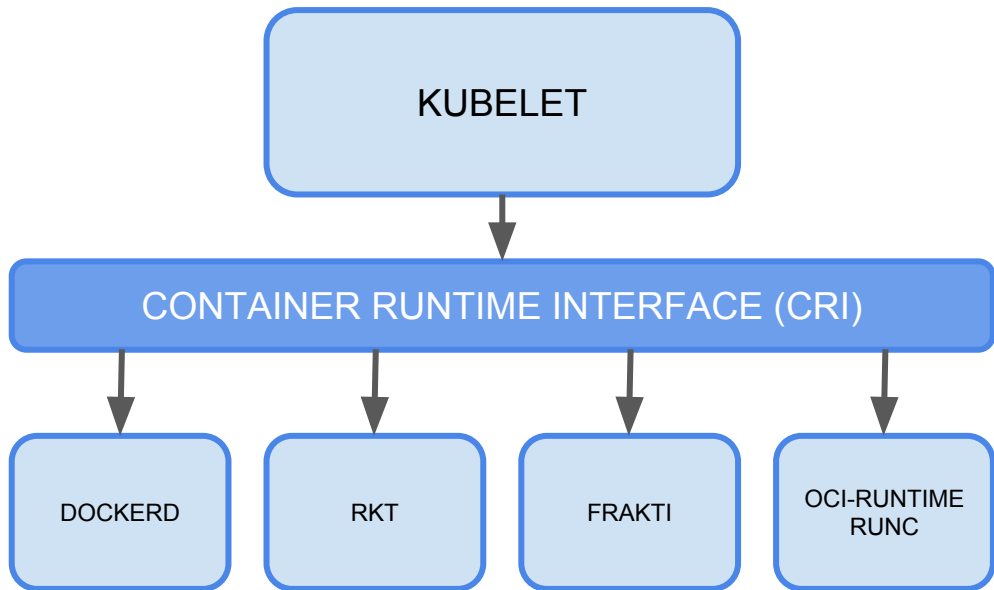



# Standardize containers



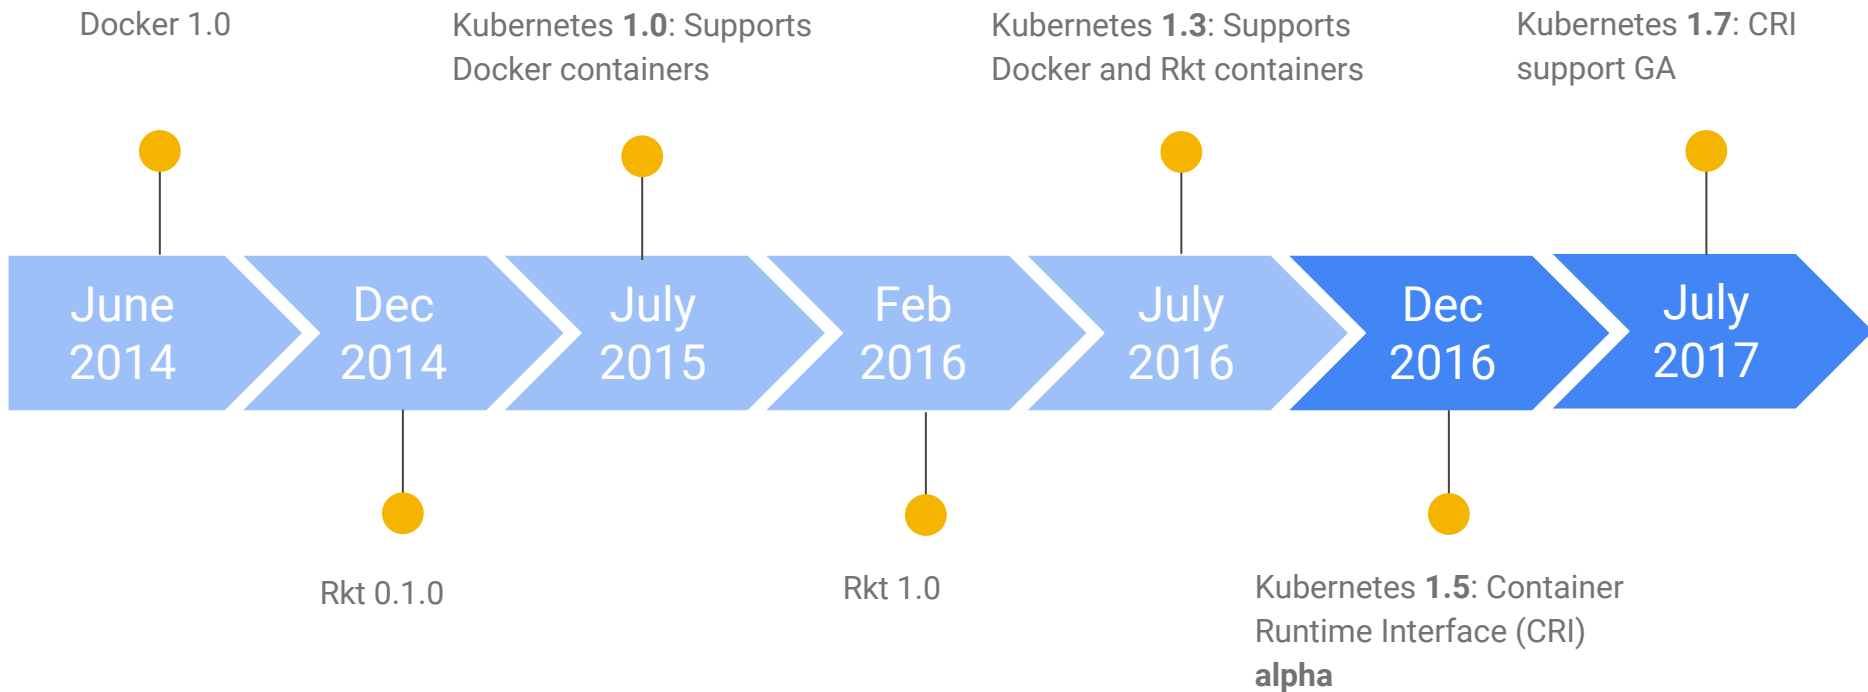
- **Runtime spec (runc = Reference implementation)**
- **Image spec**
- **Distribution spec (proposal)**

# Use API/Interfaces to Container Runtimes






# Standardization became a fact



# What is Container Runtime Interface (CRI)?

- A gRPC interface and a group of libraries
- Enables Kubernetes to use a wide variety of container runtimes
- Introduced in Kubernetes 1.5
- GA in Kubernetes 1.7



# CRI Implementations



cri-o



frakti



virtlet



dockershim

# CRI-O

- Open source & Open governance
- Lean, Stable, Secure and BORING!
  - Tied to the **CRI**
  - No features that can mine **stability** and **performance**
  - **Shaped** around Kubernetes
  - **Only supported user** is Kubernetes
  - **Versioning and Support** are tied to Kubernetes





**Man, this guy is so boring!**

**When is the live demo?**





What if I want to try it?

```
$ minikube start \  
  --network-plugin=cni \  
  --container-runtime=cri-o \  
  --bootstrapper=kubeadm
```



# skopeo

- **Copy** images from/to (multiple transports/storages):
  - containers-storage:docker-reference
  - dir:path
  - docker://docker-reference
  - docker-archive:path[:docker-reference]
  - docker-daemon:docker-reference
  - oci:path:tag
  - ostree:image[@/absolute/repo/path]
- **Inspect** images
- **Delete an** image from a repository
- **Standalone binary / No daemon running**
- Perfect for pipelines (e.g. Jenkins)





# buildah

- Build images
- **No daemon running**
- shell-like syntax
- Build from Dockerfile(s)







# libpod/podman

Library (libpod) and CLI (podman) for managing OCI-based Pods, Containers, and Container Images

- Replacement for docker cli
  - known CLI
- Integrated with CRI-O (soon)
- **No daemon running**



I wanna

SEE



“Our ancestors called it  
magic, but you call it  
[computer] science.

I come from a land where  
they are one and the same.”

—THOR

What else?



# Daemon-less Dockerfile builds

- Consume a Dockerfile, but build image without a docker daemon
- Pros
  - Docker build-like experience (just write a Dockerfile)
  - Potentially more control over image layers (combine or shard)
  - Aim is for greater security
- Cons
  - Dockerfile fidelity might make difficult some use cases
  - Different approaches to image layer construction



# Daemon-less Dockerfile builds

- [Buildah](#)
  - a tool that facilitates building OCI container images
- [Img](#)
  - Standalone, daemon-less, unprivileged Dockerfile and OCI compatible container image builder.
  - The commands/UX are the same as docker (drop-in replacement)
- [Kaniko](#)
  - kaniko is a tool to build OCI container images from a Dockerfile, inside a container or Kubernetes cluster
  - executes each command within a Dockerfile completely in userspace
- more...

# Dockerfile-less builds

- User input is source / intent: “I want to run a Node.js web server”
- Pros:
  - Less configuration
  - Tools can intelligently build layers, better/safe layer caching
  - Docker image best practices can be codified into tools
- Cons:
  - Less flexible - Opinionated builds
  - Very fragmented across vendors, no real standard

# Dockerfile-less builds

- Source to Image
  - User provides source, source gets built+layered into an application image
  - Dependent on ecosystem of framework/language builder images
- Buildpacks
  - Invented by Heroku, adopted by Cloud Foundry / Deis
  - User provides source, “build” produces “slug”, “export” produces container image
- FTL (Faster than light)
  - Purpose-built source to image builders per-language, goal is layer-per-dependency
  - Insight: turn build incrementality into deploy incrementality
- Bazel
  - Google’s OSS build system, supports declarative image builds
  - Used for user-mode Docker image builds for 3+ years

And don't forget to  
tweet if you liked it  
(or if you didn't)  
[@jorgemoralespou](#)

