



Build and run applications in a Dockerless Kubernetes world

Jorge Morales
OpenShift Developer Advocate
Riga Dev Days 2018





Me (aka Jorge Morales)



- Spanish by nature and by language
- Work at Red Hat
- OpenShift Developer Advocate
- Mostly Java developer
- Obsessed with improving the developer experience



<http://jorgemoral.es>



@jorgemoralespou



github.com/jorgemoralespou

A photograph of a dirt road winding through a forest. The road is surrounded by tall, dry grass and leafless bushes on the left, and a dense line of evergreen trees on the right. The sky is overcast.

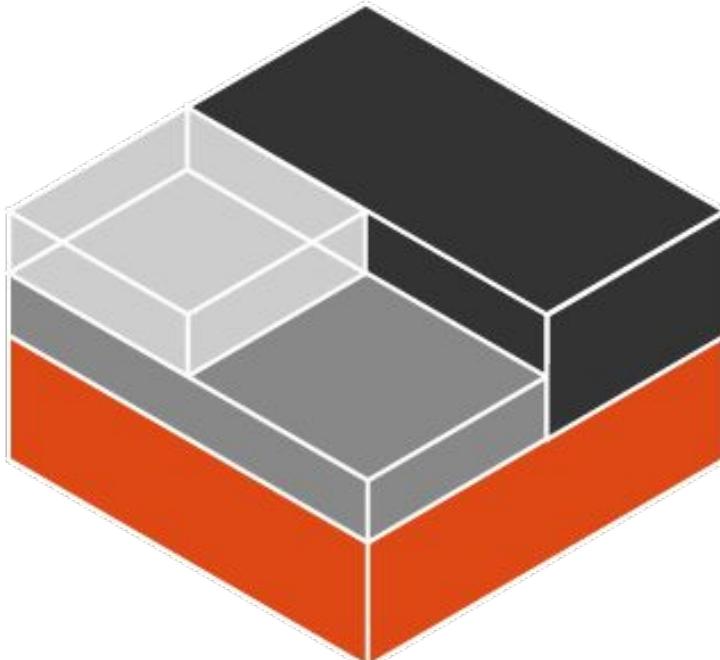
**Each day is a little
bit of history.**

José Saramago

ONCE UPON A TIME



Linux Containers



Kernel namespaces: sandboxing processes from one another

Control Groups (cgroups): control process resource allocations

Security: capabilities drop (seccomp), Mandatory access control (SELinux, Apparmor)

VMS vs Containers



VMs? Containers?
All I want to do is program!
Jeesh.





Making things easy is hard.

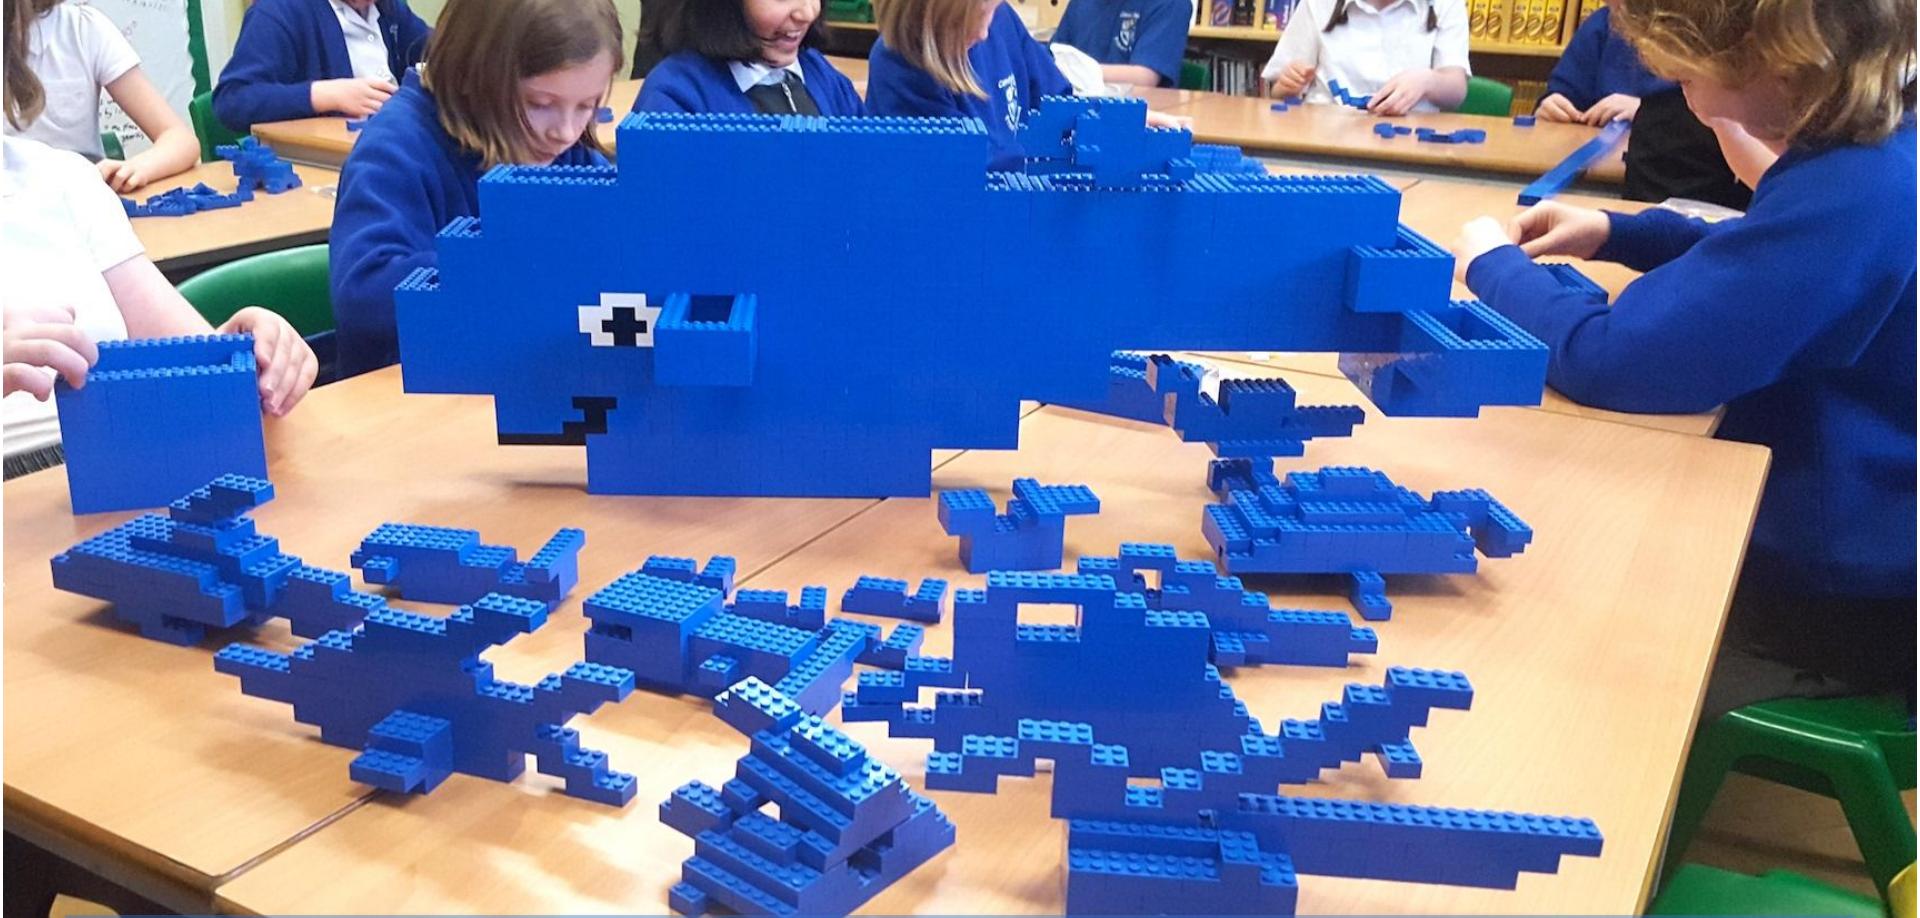
Ted Nelson



@jorgemoralespou

*Because I'm
Happy!*





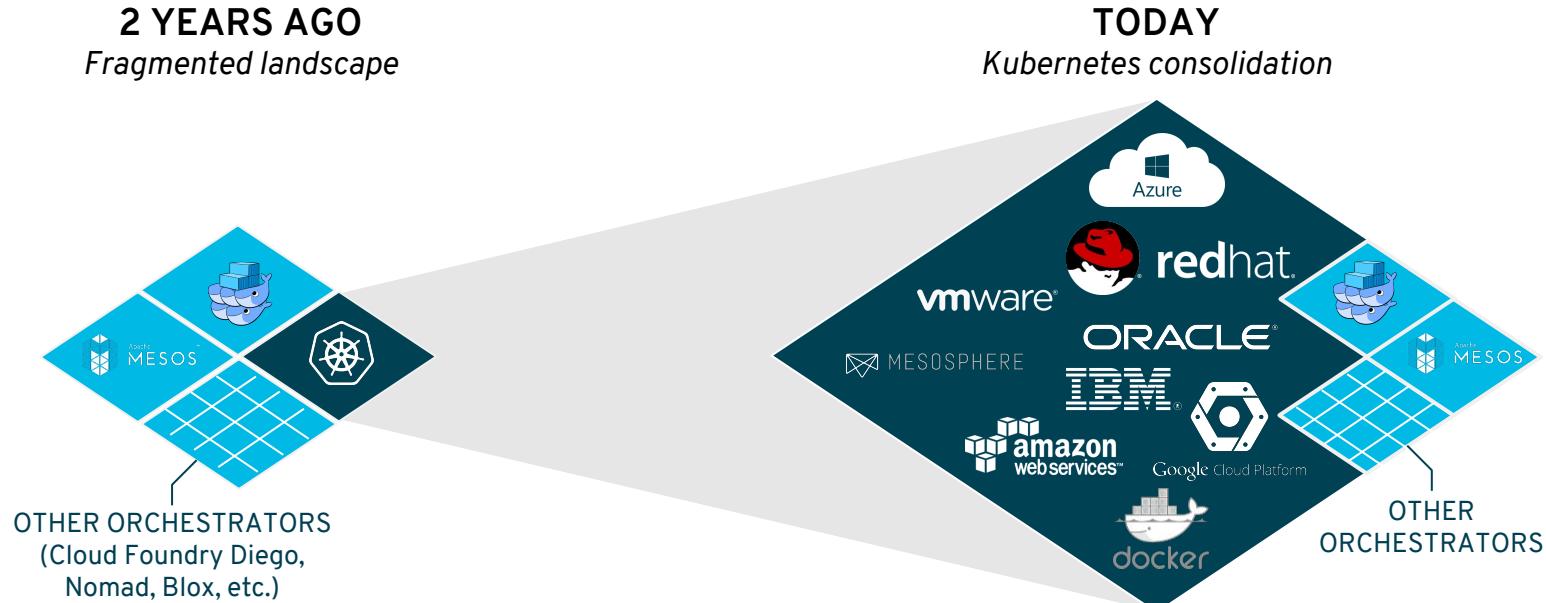
With scale came complexity



Orchestration

@jorgemoralespou

CONTAINER ORCHESTRATION LANDSCAPE



Why kubernetes?

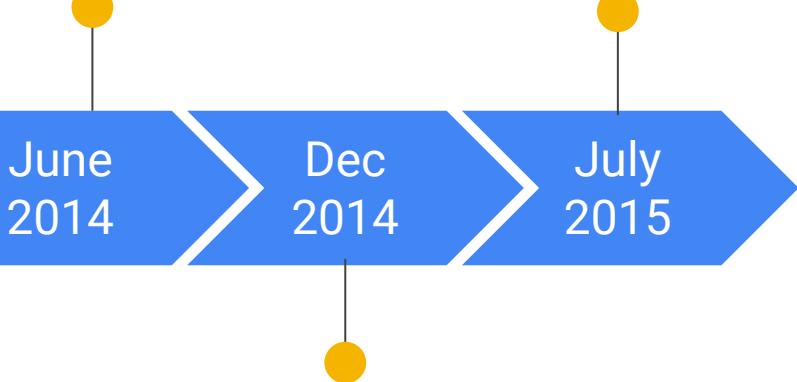


- #1: Open source, backed by giants
- #2: Vibrant and fast growing community
- #3: Supported on all clouds
- #4: Great partnerships

Started slow

Docker 1.0

Kubernetes 1.0: Supports
Docker containers



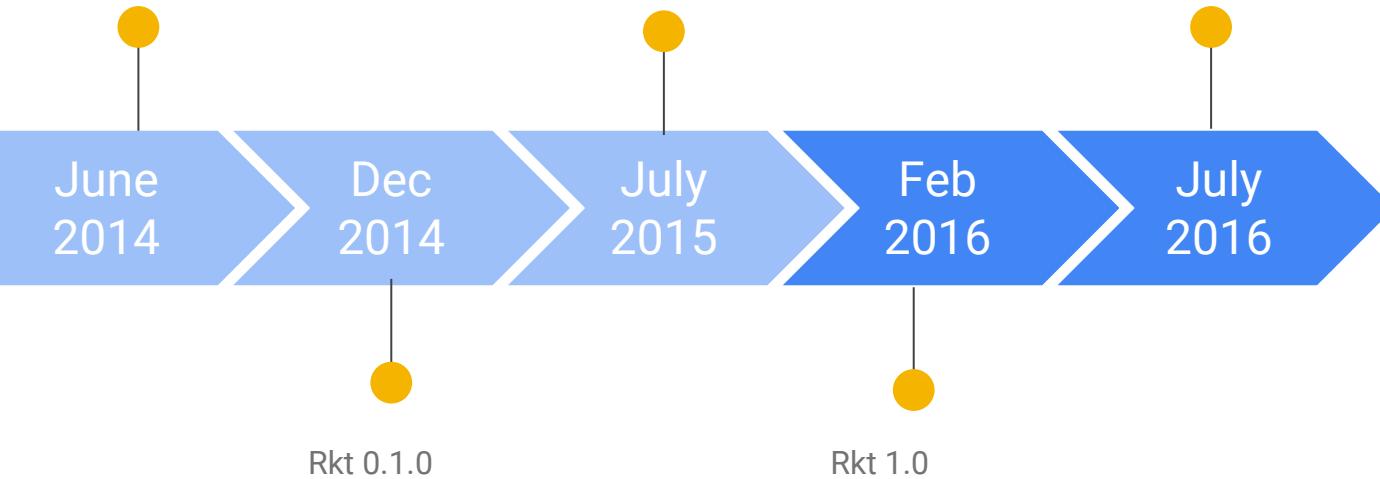
Rkt 0.1.0

then more runtimes showed up

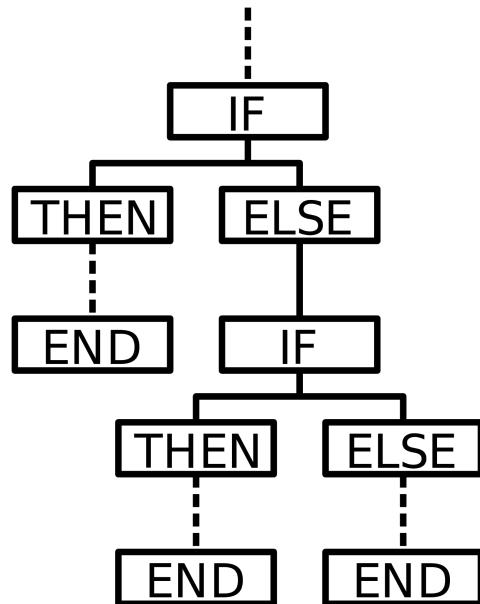
Docker 1.0

Kubernetes 1.0: Supports
Docker containers

Kubernetes 1.3: Supports
Docker and Rkt containers



and code got messy



“Change is the essential process of all of existence.”

—SPOCK

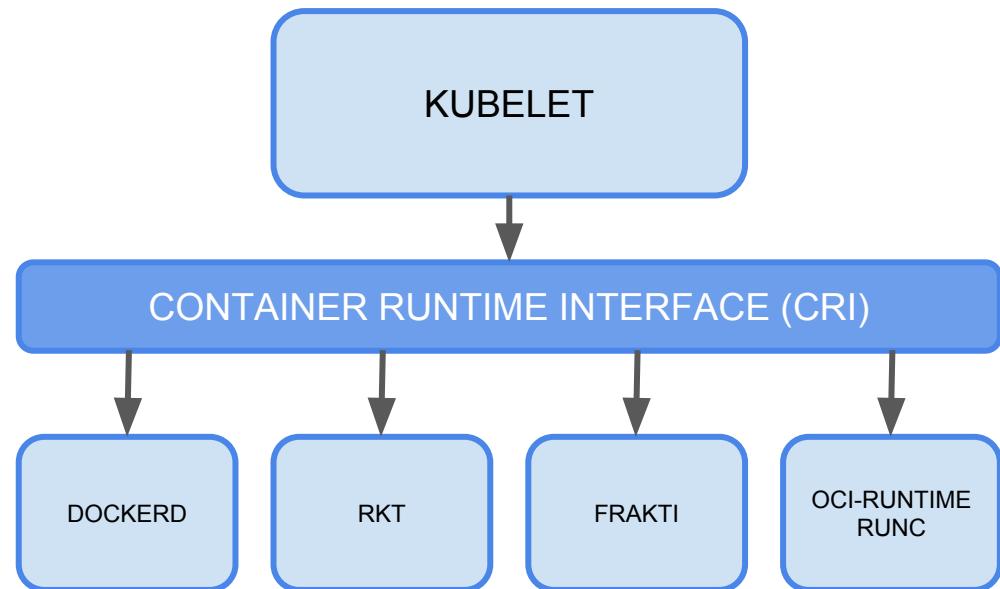
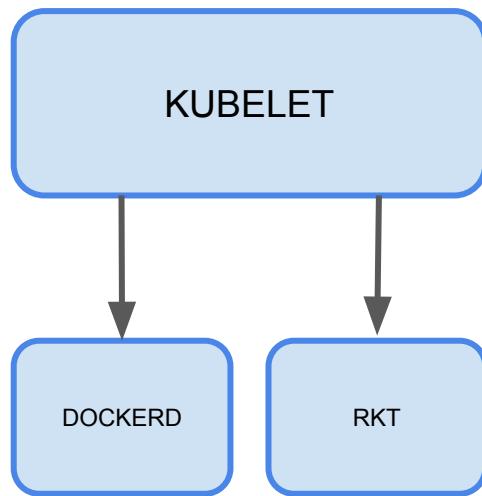


Standardize containers



- Runtime spec (`runc` = Reference implementation)
- Image spec
- Distribution spec (proposal)

Use API/Interfaces to Container Runtimes



Standardization became a fact

Docker 1.0

Kubernetes 1.0: Supports Docker containers

Kubernetes 1.3: Supports Docker and Rkt containers

Kubernetes 1.7: CRI support GA

June
2014

Dec
2014

July
2015

Feb
2016

July
2016

Dec
2016

July
2017

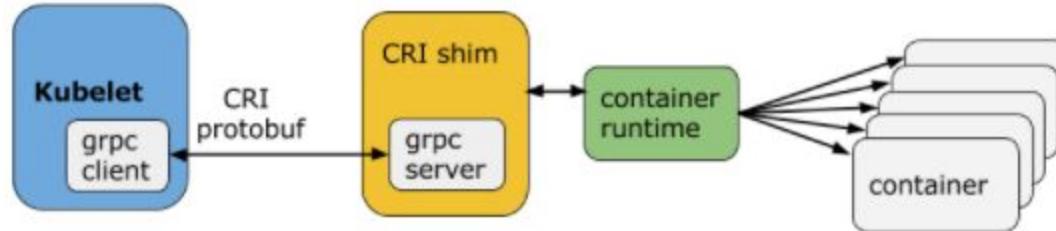
Rkt 0.1.0

Rkt 1.0

Kubernetes 1.5: Container Runtime Interface (CRI)
alpha

What is Container Runtime Interface (CRI)?

- A gRPC interface and a group of libraries
- Enables Kubernetes to use a wide variety of container runtimes
- Introduced in Kubernetes 1.5
- GA in Kubernetes 1.7



CRI Implementations



cri-o

containerd
cri-containerd



frakti



rktlet



virtlet

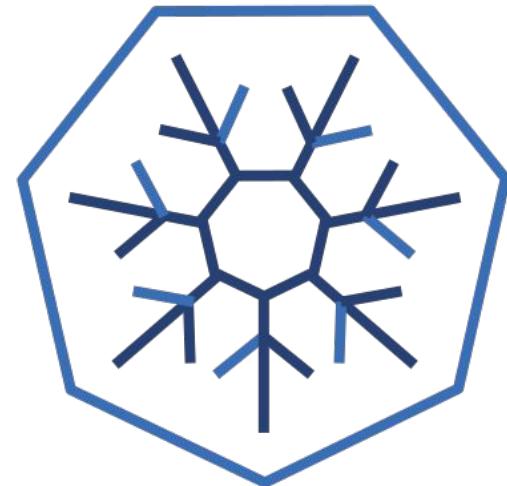


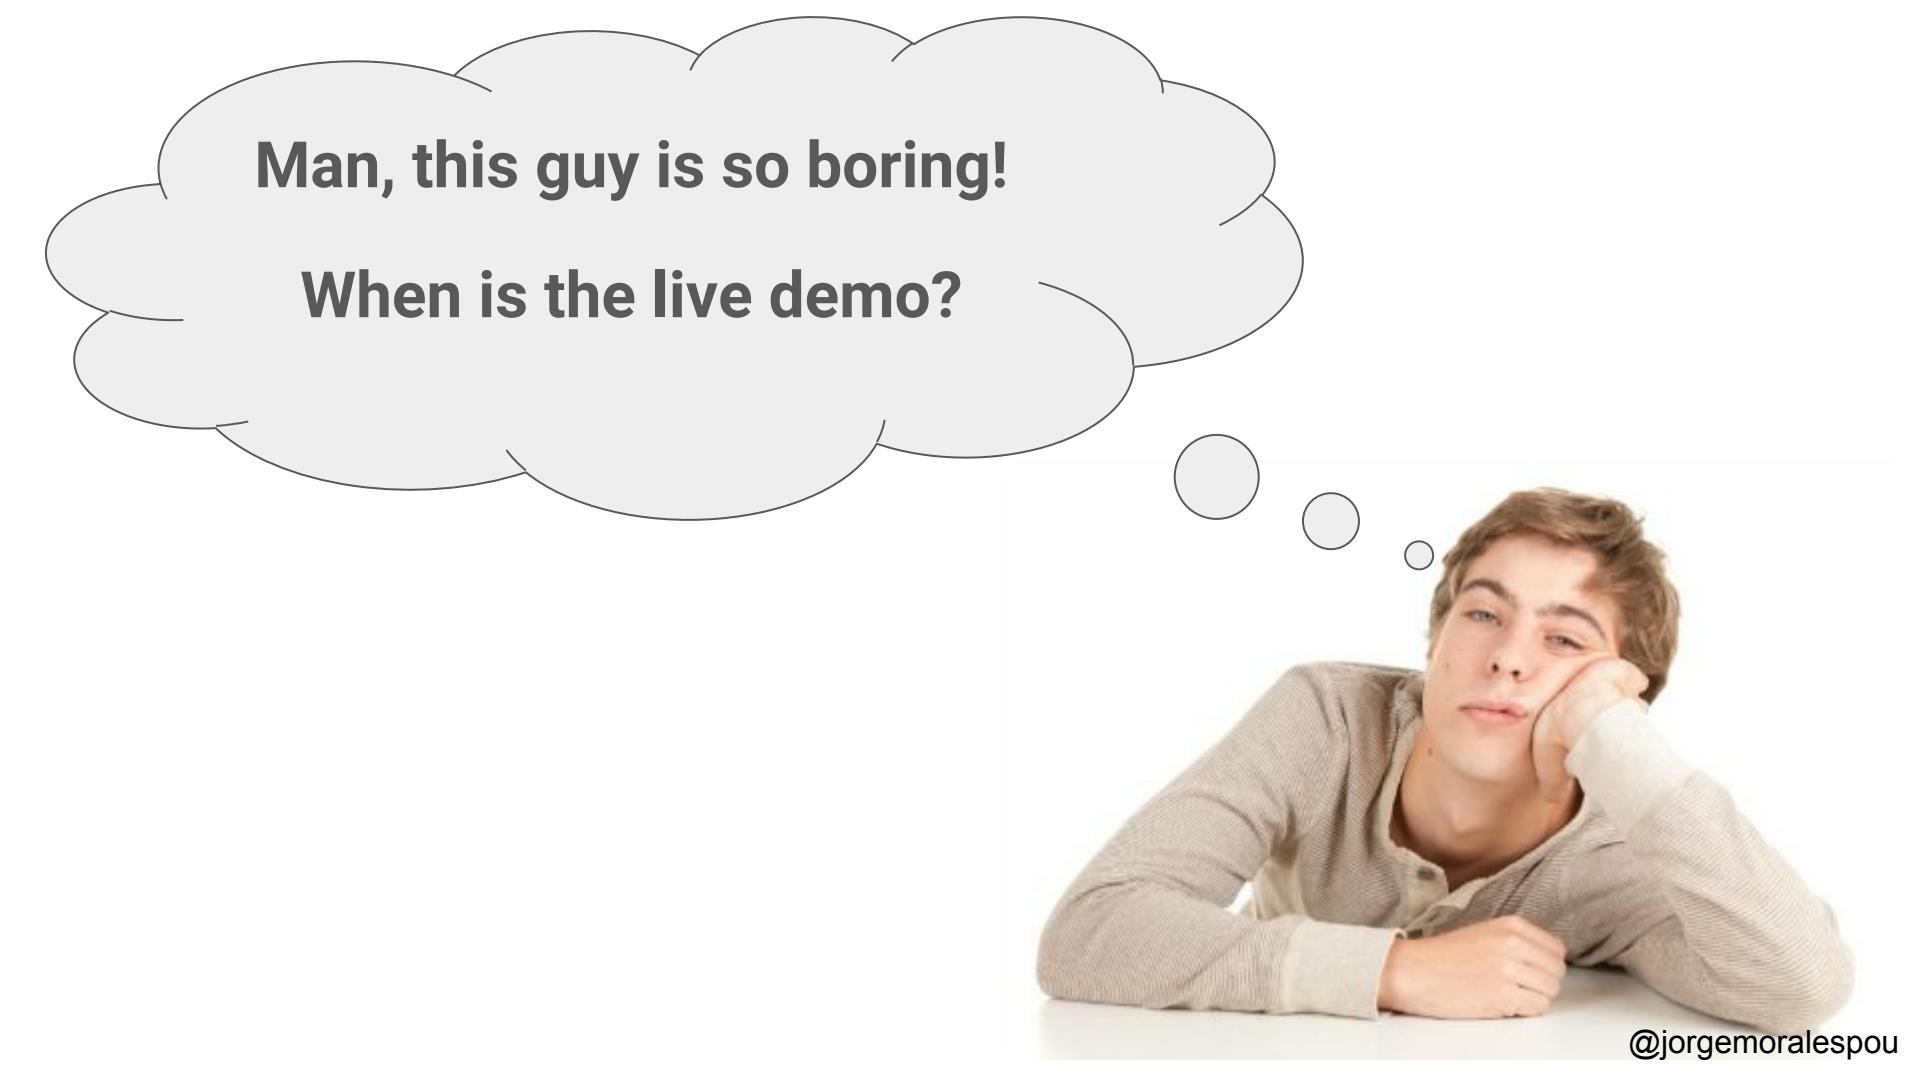
dockershim

@jorgemoralespou

CRI-O

- Open source & Open governance
- Lean, Stable, Secure and BORING!
 - Tied to the **CRI**
 - No features that can mine **stability** and **performance**
 - **Shaped** around Kubernetes
 - **Only supported user** is Kubernetes
 - **Versioning and Support** are tied to Kubernetes





Man, this guy is so boring!

When is the live demo?

Demo script

SHOW DOCKER AND STOP IT

```
$ docker images  
$ systemctl stop docker  
$ docker images
```

RUN A CONTAINER WITH OC/KUBECTL

```
$ kubectl get pods  
$ kubectl run --image=nginx --port=80 nginx  
$ kubectl expose deployment nginx --port=80  
$ kubectl get svc  
$ curl http://<CLUSTER-IP>  
    $ oc expose svc/nginx-http  
    $ oc get route  
$ kubectl get pods  
$ kubectl logs -f <POD_NAME>  
$ kubectl exec -it <POD_NAME> sh  
$ docker ps  
$ runc list (show that containers are run by runc)
```



What if I want to try it?

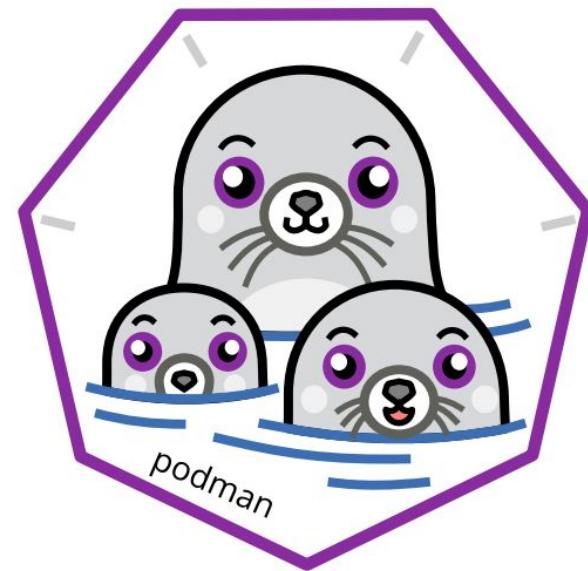
```
$ minikube start \  
    --network-plugin=cni \  
    --container-runtime=cri-o \  
    --bootstrapper=kubeadm
```



libpod/podman

Library (libpod) and CLI (podman) for managing OCI-based Pods, Containers, and Container Images

- Replacement for docker cli
 - known CLI
- Integrated with CRI-O (soon)
- **No daemon running**



I WANNA
SEE

Demo script

```
$ podman ps
$ podman images
$ podman run -it --rm -p 8080:80 nginx
$ <ANOTHER TERMINAL> curl http://localhost:8080
$ <ANOTHER TERMINAL> podman ps
$ <ANOTHER TERMINAL> podman logs <CONTAINER_ID>
$ <ANOTHER TERMINAL> podman exec -t <CONTAINER_ID> sh
$ <ANOTHER TERMINAL> ls /etc/nginx

$ podman images
$ podman run -t fedora echo "Hello Riga"
$ podman ps -a

$ cd Dockerfiles
$ podman build -t rigacontainer .
$ podman images
$ podman run -it rigacontainer cat /hello
$ podman ps -a
$ podman rm --all
```



skopeo

- **Copy** images from/to (multiple transports/storages):
 - containers-storage:docker-reference
 - dir:path
 - docker://docker-reference
 - docker-archive:path[:docker-reference]
 - docker-daemon:docker-reference
 - oci:path:tag
 - ostree:image[@/absolute/repo/path]
- **Inspect** images
- **Delete** an image from a repository
- **Standalone binary / No daemon running**
- Perfect for pipelines (e.g. Jenkins)



I WANNA
SEE

Demo script

```
$ systemctl start docker  
$ docker images  
$ podman images
```

COPY IMAGE

```
$ skopeo copy containers-storage:riga/example:latest docker-daemon:riga/example:latest  
$ docker images  
$ docker run -it --rm riga/example cat /hello  
$ docker ps -a  
$ docker rm $(docker ps -qa)  
$ skopeo inspect docker-daemon:riga/example:latest  
$ skopeo inspect docker:docker.io/library/fedora:latest
```



buildah

- Build images
- **No daemon running**
- shell-like syntax
- Build from Dockerfile(s)



I WANNA
SEE

Demo script

```
$ cd ~/Dockerfiles
$ skopeo copy containers-storage:registry.fedoraproject.org/fedora:latest docker-daemon:fedora:latest
$ docker build -t riga/example-docker .

$ docker history riga/example
$ docker history riga/example-docker

$ buildah bud -t riga/buildah-dockerfile

$ container=(buildah from fedora)
$ echo $container
$ buildah containers
$ buildah config --author "Jorge" --label "METADATA=Built with buildah" $container
$ buildah inspect $container
$ buildah run $container bash
# echo "Hello Riga, built by Buildah" > /hello
# ls /
# cat /hello
# exit
$ buildah commit $container riga/example-buildah
$ podman run -it riga/example-buildah cat /hello
```

What else?



@jorgemoralespou

Daemon-less Dockerfile builds

- Consume a Dockerfile, but build image without a docker daemon
- Pros
 - Docker build-like experience (just write a Dockerfile)
 - Potentially more control over image layers (combine or shard)
 - Aim is for greater security
- Cons
 - Dockerfile fidelity might make difficult some use cases
 - Different approaches to image layer construction

Daemon-less Dockerfile builds

- [Buildah](#)
 - a tool that facilitates building OCI container images
- [Img](#)
 - Standalone, daemon-less, unprivileged Dockerfile and OCI compatible container image builder.
 - The commands/UX are the same as docker (drop-in replacement)
- [Kaniko](#)
 - kaniko is a tool to build OCI container images from a Dockerfile, inside a container or Kubernetes cluster
 - executes each command within a Dockerfile completely in userspace
- more...

Dockerfile-less builds

- User input is source / intent: “I want to run a Node.js web server”
- Pros:
 - Less configuration
 - Tools can intelligently build layers, better/safe layer caching
 - Docker image best practices can be codified into tools
- Cons:
 - Less flexible - Opinionated builds
 - Very fragmented across vendors, no real standard

Dockerfile-less builds

- Source to Image
 - User provides source, source gets built+layered into an application image
 - Dependent on ecosystem of framework/language builder images
- Buildpacks
 - Invented by Heroku, adopted by Cloud Foundry / Deis
 - User provides source, “build” produces “slug”, “export” produces container image
- FTL (Faster than light)
 - Purpose-built source to image builders per-language, goal is layer-per-dependency
 - Insight: turn build incrementality into deploy incrementality
- Bazel
 - Google’s OSS build system, supports declarative image builds
 - Used for user-mode Docker image builds for 3+ years



“Our ancestors called it
magic, but you call it
[computer] science.

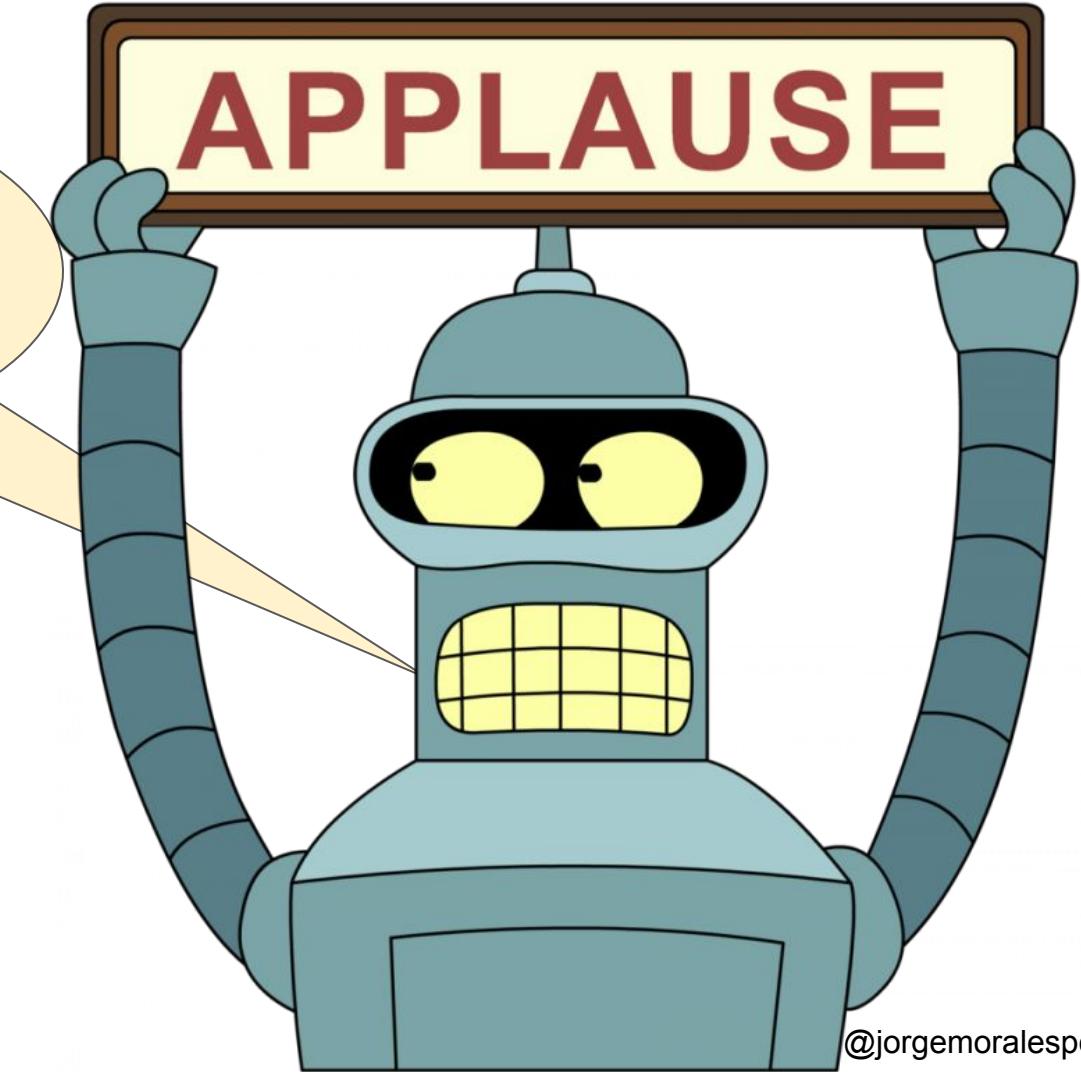
I come from a land where
they are one and the same.”

—THOR



@jorgemoralespou

And don't forget to
tweet if you liked it
(or if you didn't)
@jorgemoralespou



@jorgemoralespou